

190020066_assignment

April 21, 2021

1 Programming Assignment

- Author: Laxman Desai (190020066)
- Course: CL 202 (Data Analysis)
- Date: 20th April 2021

1.1 Initialization

Z_1 = Total dwelling space (100 ft^2)

Z_2 = Assessed value (\$1000)

Y = Selling price (\$1000)

```
[2]: import pandas as pd
import numpy as np

df = pd.read_csv('dataset/data.csv')
df['Z1'] *= 100 # Scaling
df['Z2'] *= 1000 # Scaling
df['Y'] *= 1000 # Scaling

df.head()
```

```
[2]:      Z1      Z2      Y
0  1531.0  57300.0  74800.0
1  1520.0  63800.0  74000.0
2  1625.0  65400.0  72900.0
3  1433.0  57000.0  70000.0
4  1457.0  63800.0  74900.0
```

1.2 A) Correlation Coefficient

```
[3]: display(df.corr())

corr_coeff_Y_Z1 = df['Y'].corr(df['Z1'])
print(f'corr_coeff_Y_Z1 = {round(corr_coeff_Y_Z1, 4)}')

corr_coeff_Y_Z2 = df['Y'].corr(df['Z2'])
print(f'corr_coeff_Y_Z2 = {round(corr_coeff_Y_Z2, 4)}')
```

	Z1	Z2	Y
Z1	1.000000	0.925711	0.913319
Z2	0.925711	1.000000	0.851384
Y	0.913319	0.851384	1.000000

corr_coeff_Y_Z1 = 0.9133

corr_coeff_Y_Z2 = 0.8514

1.3 B) Linear Model

$$Y = a*Z1 + b*Z2 + c$$

```
[4]: X = df.drop('Y', axis=1) # Independent Variables
      y = df['Y'] # Dependant Variables

      N = len(X)
      print(f'No. of training samples = {N}')

      p = len(X.columns) + 1 # '+1' because LinearRegression adds an intercept term
      print(f'No. of regression coefficients = {p}, i.e. Y = a*Z1 + b*Z2 + c*1')
```

No. of training samples = 20

No. of regression coefficients = 3, i.e. $Y = a*Z1 + b*Z2 + c*1$

```
[5]: from sklearn.linear_model import LinearRegression

      model = LinearRegression()
      model.fit(X, y)

      a, b = model.coef_
      c = model.intercept_

      print(f'a = {round(a, 5)}, b = {round(b, 5)}, c = {round(c, 5)}')
```

a = 26.344, b = 0.04518, c = 30966.56634

1.4 C) 95% Confidence Interval

$$ci_l_a < a < ci_u_a$$

$$ci_l_b < b < ci_u_b$$

$$ci_l_c < c < ci_u_c$$

```
[6]: y_hat = model.predict(X)
      err = y - y_hat # residual

      RSS = err.T @ err # dot product, equivalent to sum(err**2)
      print(f'Residual sum of squares = {round(RSS, 6)}')

      S = np.sqrt(RSS / (N - p))
      print(f'Standard deviation of y_hat = {S}')
```

```
Residual sum of squares = 204994944.947883
Standard deviation of y_hat = 3472.5388656435357
```

```
[7]: # M = matrix (20 x 3) with columns as [Z1, Z2, 1]
M = np.ones(shape=(N, p), dtype=float)
M[:, 0] = X.iloc[:, 0]
M[:, 1] = X.iloc[:, 1]

display(M[0:5]) # Equivalent to M.head() for non dataframes
```

```
array([[1.531e+03, 5.730e+04, 1.000e+00],
       [1.520e+03, 6.380e+04, 1.000e+00],
       [1.625e+03, 6.540e+04, 1.000e+00],
       [1.433e+03, 5.700e+04, 1.000e+00],
       [1.457e+03, 6.380e+04, 1.000e+00]])
```

```
[8]: var_beta_hat = 1/(M.T @ M) * S**2 # Variance of beta hat

std_err_a = var_beta_hat[0, 0] ** 0.5
print(f'Standard Error of a = {std_err_a}')

std_err_b = var_beta_hat[1, 1] ** 0.5
print(f'Standard Error of b = {std_err_b}')

std_err_c = var_beta_hat[2, 2] ** 0.5
print(f'Standard Error of c = {std_err_c}')
```

```
Standard Error of a = 0.47258255944417654
Standard Error of b = 0.012232992294284532
Standard Error of c = 776.4832958088955
```

```
[13]: # from statistics import NormalDist
# z = NormalDist().inv_cdf((1 + 0.95) / 2.) # 95% CI
# print(f'z for 95% CI = {z}')
```

```
from scipy import stats

t = abs(stats.t.ppf((1-0.95)/2, N-p)) # 95% CI
print(f't for 95% CI = {t}')
```

```
t for 95% CI = 2.1098155778331806
```

```
[14]: # 95% CI Lower Bounds
ci_l_a = a - t * std_err_a
ci_l_b = b - t * std_err_b
ci_l_c = c - t * std_err_c

# 95% CI Upper Bounds
ci_u_a = a + t * std_err_a
```

```

ci_u_b = b + t * std_err_b
ci_u_c = c + t * std_err_c

print(f'a    [{round(ci_l_a, 6)}, {round(ci_u_a, 6)}]')
print(f'b    [{round(ci_l_b, 6)}, {round(ci_u_b, 6)}]')
print(f'c    [{round(ci_l_c, 6)}, {round(ci_u_c, 6)}]')

```

```

a    [25.346934, 27.341058]
b    [0.019375, 0.070993]
c    [29328.329782, 32604.802889]

```

1.4.1 Plotting the data:

```

[15]: # Plotting the linear model
      # %matplotlib ipympl
      %matplotlib inline
      import matplotlib.pyplot as plt

      # Initialize matplotlib
      plt.style.use('ggplot')
      plt.rc('text', usetex=True) # Unnecessary
      plt.rcParams['text.latex.preamble'] = r'\usepackage{amsmath}'
      plt.rcParams['grid.color'] = '#C0C0C0'

      # Creating figures
      fig = plt.figure(figsize=(16, 6))
      fig.set_facecolor('#FFFFFF')

      ax1 = fig.add_subplot(131, projection='3d')
      ax2 = fig.add_subplot(132, projection='3d')
      ax3 = fig.add_subplot(133, projection='3d')
      ax1.view_init(elev=28, azim=120)
      ax2.view_init(elev=5, azim=-80)
      ax3.view_init(elev=60, azim=165)
      axes = [ax1, ax2, ax3]

      # Create regression grid
      x_plot = np.linspace(1300, 2600, 30) # range of Z1 to plot
      y_plot = np.linspace(50000, 90000, 30) # range of Z2 to plot
      x_mesh, y_mesh = np.meshgrid(x_plot, y_plot)

      model_viz = np.array([x_mesh.flatten(), y_mesh.flatten()]).T
      y_hat_plot = model.predict(model_viz)

      # Creating plot
      for ax in axes:

```

```

ax.plot(df['Z1'], df['Z2'], df['Y'], color='k', zorder=15,
↳linestyle='none', marker='o', alpha=0.5)
ax.scatter(x_mesh, y_mesh, y_hat_plot, s=60, facecolor='#1F77B422',
↳edgecolor='#1F77B4cc')

ax.set_xlabel(r'$\boldsymbol{Z_1 \ (ft^2)}$', fontsize=12)
ax.set_ylabel(r'$\boldsymbol{Z_2 \ (\$)}$', fontsize=12)
ax.set_zlabel(r'$\boldsymbol{Y \ (\$)}$', fontsize=12)


ax.set_facecolor('#FFFFFF')
ax.locator_params(nbins=4, axis='x')
ax.locator_params(nbins=4, axis='x')
ax.ticklabel_format(style="sci", scilimits=(0, 0), useMathText=True)

fig.suptitle('Linear Regression Plot (3D)', fontsize=30)
fig.tight_layout(pad=1.5)

# plt.savefig('Linear Regression Plot.png')

```

Linear Regression Plot (3D)



1.5 D) 95% Prediction Interval

- Apartment size = 12 * 100 ft² and assessed value of 60 * 1000\$
- $\pi_{l_Y} < \pi_u < \pi_{u_Y}$

```
[16]: X_test = [[1200, 60000]]
```

```

y_test_hat = model.predict(X_test)[0]
print(f'Predicted selling price = {round(y_test_hat, 2)}$')

```

Predicted selling price = 65290.39\$

```
[17]: pi_l_Y = y_test_hat - t * S
      pi_u_Y = y_test_hat + t * S
      print(f'lower value bound = {pi_l_Y}, upper value bound = {pi_u_Y}')
```

lower value bound = 57963.9768584323, upper value bound = 72616.8100451641

1.6 E) Mean & variance of residuals

```
[18]: mean_resid = err.mean()
      print(f'mean of residual = {mean_resid}')

      var_resid = err.var()
      print(f'variance of residual = {var_resid}')
```

mean of residual = 4.365574568510055e-12
variance of residual = 10789207.628835956

1.7 F) R^2 of the fit

```
[19]: y_diff = y - y.mean()
      TSS = y_diff.T @ y_diff

      Rsq = 1 - RSS/TSS
      print(f'R^2 of the fit (calculated) = {Rsq}')
```

R^2 of the fit (calculated) = 0.8343970328484549

```
[20]: # Confirming using sklearn's builtin method
      from sklearn.metrics import r2_score
      print(f'R^2 of the fit (using builtin methods) = {r2_score(y, y_hat)}')
```

R^2 of the fit (using builtin methods) = 0.8343970328484549

References: [1](#) [2](#) [3](#) [4](#)